

## breastcancer

January 5, 2026

```
[94]: import pandas as pd
import numpy as np
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix,
                           roc_auc_score, roc_curve, accuracy_score, recall_score
```

```
[106]: df = pd.read_csv(r"C:\Users\karan\Downloads\Breast_Cancer_data.csv")
```

```
[108]: df
```

```
[108]:      id diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean \
0     842302        M       17.99       10.38       122.80      1001.0
1     842517        M       20.57       17.77       132.90      1326.0
2    84300903        M       19.69       21.25       130.00      1203.0
3    84348301        M       11.42       20.38        77.58      386.1
4    84358402        M       20.29       14.34       135.10      1297.0
..      ...
564   926424        M       21.56       22.39       142.00      1479.0
565   926682        M       20.13       28.25       131.20      1261.0
566   926954        M       16.60       28.08       108.30      858.1
567   927241        M       20.60       29.33       140.10      1265.0
568   92751         B        7.76       24.54        47.92      181.0

      smoothness_mean  compactness_mean  concavity_mean  concave points_mean \
0            0.11840           0.27760          0.30010           0.14710
1            0.08474           0.07864          0.08690           0.07017
2            0.10960           0.15990          0.19740           0.12790
3            0.14250           0.28390          0.24140           0.10520
```

4	0.10030	0.13280	0.19800	0.10430
..	...	...	...	...
564	0.11100	0.11590	0.24390	0.13890
565	0.09780	0.10340	0.14400	0.09791
566	0.08455	0.10230	0.09251	0.05302
567	0.11780	0.27700	0.35140	0.15200
568	0.05263	0.04362	0.00000	0.00000
	... texture_worst	perimeter_worst	area_worst	smoothness_worst \
0	... 17.33	184.60	2019.0	0.16220
1	... 23.41	158.80	1956.0	0.12380
2	... 25.53	152.50	1709.0	0.14440
3	... 26.50	98.87	567.7	0.20980
4	... 16.67	152.20	1575.0	0.13740
..	... ...	...	...	...
564	... 26.40	166.10	2027.0	0.14100
565	... 38.25	155.00	1731.0	0.11660
566	... 34.12	126.70	1124.0	0.11390
567	... 39.42	184.60	1821.0	0.16500
568	... 30.37	59.16	268.6	0.08996
	compactness_worst	concavity_worst	concave points_worst	symmetry_worst \
0	0.66560	0.7119	0.2654	0.4601
1	0.18660	0.2416	0.1860	0.2750
2	0.42450	0.4504	0.2430	0.3613
3	0.86630	0.6869	0.2575	0.6638
4	0.20500	0.4000	0.1625	0.2364
..	... ...	...	...	...
564	0.21130	0.4107	0.2216	0.2060
565	0.19220	0.3215	0.1628	0.2572
566	0.30940	0.3403	0.1418	0.2218
567	0.86810	0.9387	0.2650	0.4087
568	0.06444	0.0000	0.0000	0.2871
	fractal_dimension_worst	Unnamed: 32		
0	0.11890	NaN		
1	0.08902	NaN		
2	0.08758	NaN		
3	0.17300	NaN		
4	0.07678	NaN		
..	... ...	...		
564	0.07115	NaN		
565	0.06637	NaN		
566	0.07820	NaN		
567	0.12400	NaN		
568	0.07039	NaN		

```
[569 rows x 33 columns]
```

```
[110]: skewness_values = df.skew(numeric_only=True)
print(skewness_values)
```

```
id                6.473752
radius_mean       0.942380
texture_mean      0.650450
perimeter_mean    0.990650
area_mean         1.645732
smoothness_mean   0.456324
compactness_mean  1.190123
concavity_mean    1.401180
concave_points_mean 1.171180
symmetry_mean     0.725609
fractal_dimension_mean 1.304489
radius_se          3.088612
texture_se          1.646444
perimeter_se        3.443615
area_se             5.447186
smoothness_se       2.314450
compactness_se      1.902221
concavity_se         5.110463
concave_points_se   1.444678
symmetry_se          2.195133
fractal_dimension_se 3.923969
radius_worst        1.103115
texture_worst        0.498321
perimeter_worst      1.128164
area_worst           1.859373
smoothness_worst     0.415426
compactness_worst    1.473555
concavity_worst       1.150237
concave_points_worst 0.492616
symmetry_worst        1.433928
fractal_dimension_worst 1.662579
Unnamed: 32            NaN
dtype: float64
```

```
[112]: df.shape
```

```
[112]: (569, 33)
```

```
[114]: df.isnull().sum()
```

```
[114]: id              0
diagnosis          0
radius_mean         0
```

```

texture_mean          0
perimeter_mean        0
area_mean             0
smoothness_mean       0
compactness_mean      0
concavity_mean        0
concave points_mean   0
symmetry_mean         0
fractal_dimension_mean 0
radius_se              0
texture_se              0
perimeter_se            0
area_se                0
smoothness_se           0
compactness_se          0
concavity_se             0
concave points_se        0
symmetry_se              0
fractal_dimension_se     0
radius_worst            0
texture_worst            0
perimeter_worst          0
area_worst               0
smoothness_worst          0
compactness_worst         0
concavity_worst            0
concave points_worst       0
symmetry_worst             0
fractal_dimension_worst     0
Unnamed: 32                  569
dtype: int64

```

[116]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
 #   Column           Non-Null Count  Dtype  
 ____ _--_
 0   id               569 non-null    int64  
 1   diagnosis        569 non-null    object 
 2   radius_mean       569 non-null    float64
 3   texture_mean      569 non-null    float64
 4   perimeter_mean    569 non-null    float64
 5   area_mean         569 non-null    float64
 6   smoothness_mean   569 non-null    float64
 7   compactness_mean  569 non-null    float64
 8   concavity_mean    569 non-null    float64

```

```

9   concave points_mean      569 non-null    float64
10  symmetry_mean           569 non-null    float64
11  fractal_dimension_mean  569 non-null    float64
12  radius_se                569 non-null    float64
13  texture_se               569 non-null    float64
14  perimeter_se             569 non-null    float64
15  area_se                  569 non-null    float64
16  smoothness_se            569 non-null    float64
17  compactness_se           569 non-null    float64
18  concavity_se             569 non-null    float64
19  concave points_se        569 non-null    float64
20  symmetry_se              569 non-null    float64
21  fractal_dimension_se     569 non-null    float64
22  radius_worst              569 non-null    float64
23  texture_worst             569 non-null    float64
24  perimeter_worst           569 non-null    float64
25  area_worst                569 non-null    float64
26  smoothness_worst          569 non-null    float64
27  compactness_worst         569 non-null    float64
28  concavity_worst           569 non-null    float64
29  concave points_worst      569 non-null    float64
30  symmetry_worst            569 non-null    float64
31  fractal_dimension_worst   569 non-null    float64
32  Unnamed: 32                 0 non-null     float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB

```

[69]: df.describe().T

	count	mean	std	min	\
radius_mean	569.0	14.062916	3.340025	6.981000	
texture_mean	569.0	19.254736	4.187510	9.710000	
perimeter_mean	569.0	91.543787	23.047218	43.790000	
area_mean	569.0	639.765202	305.343508	143.500000	
smoothness_mean	569.0	0.096266	0.013685	0.057975	
compactness_mean	569.0	0.103222	0.049386	0.019380	
concavity_mean	569.0	0.086937	0.073900	0.000000	
concave points_mean	569.0	0.048552	0.037633	0.000000	
symmetry_mean	569.0	0.180734	0.026067	0.111200	
fractal_dimension_mean	569.0	0.062604	0.006418	0.049960	
radius_se	569.0	0.384698	0.203612	0.111500	
texture_se	569.0	1.198057	0.485500	0.360200	
perimeter_se	569.0	2.699075	1.402982	0.757000	
area_se	569.0	34.959487	24.294515	6.802000	
smoothness_se	569.0	0.006876	0.002410	0.001713	
compactness_se	569.0	0.024561	0.014947	0.002252	
concavity_se	569.0	0.030038	0.020577	0.000000	

concave points_se	569.0	0.011601	0.005486	0.000000
symmetry_se	569.0	0.020047	0.006572	0.007882
fractal_dimension_se	569.0	0.003591	0.001780	0.000895
radius_worst	569.0	16.183882	4.587249	7.930000
texture_worst	569.0	25.648453	6.054406	12.020000
perimeter_worst	569.0	106.705369	31.957777	50.410000
area_worst	569.0	849.907821	475.645240	185.200000
smoothness_worst	569.0	0.132209	0.022320	0.072500
compactness_worst	569.0	0.249883	0.142851	0.027290
concavity_worst	569.0	0.268754	0.197461	0.000000
concave points_worst	569.0	0.114606	0.065732	0.000000
symmetry_worst	569.0	0.287616	0.053868	0.156500
fractal_dimension_worst	569.0	0.083342	0.015993	0.055040
		25%	50%	75%
radius_mean	11.700000	13.370000	15.780000	21.900000
texture_mean	16.170000	18.840000	21.800000	30.245000
perimeter_mean	75.170000	86.240000	104.100000	147.495000
area_mean	420.300000	551.100000	782.700000	1326.300000
smoothness_mean	0.086370	0.095870	0.105300	0.133695
compactness_mean	0.064920	0.092630	0.130400	0.228620
concavity_mean	0.029560	0.061540	0.130700	0.282410
concave points_mean	0.020310	0.033500	0.074000	0.154535
symmetry_mean	0.161900	0.179200	0.195700	0.246400
fractal_dimension_mean	0.057700	0.061540	0.066120	0.078750
radius_se	0.232400	0.324200	0.478900	0.848650
texture_se	0.833900	1.108000	1.474000	2.434150
perimeter_se	1.606000	2.287000	3.357000	5.983500
area_se	17.850000	24.530000	45.190000	86.200000
smoothness_se	0.005169	0.006380	0.008146	0.012612
compactness_se	0.013080	0.020450	0.032450	0.061505
concavity_se	0.015090	0.025890	0.042050	0.082490
concave points_se	0.007638	0.010930	0.014710	0.025318
symmetry_se	0.015160	0.018730	0.023480	0.035960
fractal_dimension_se	0.002248	0.003187	0.004558	0.008023
radius_worst	13.010000	14.970000	18.790000	27.460000
texture_worst	21.080000	25.410000	29.720000	42.680000
perimeter_worst	84.110000	97.660000	125.400000	187.335000
area_worst	515.300000	686.500000	1084.000000	1937.050000
smoothness_worst	0.116600	0.131300	0.146000	0.190100
compactness_worst	0.147200	0.211900	0.339100	0.626950
concavity_worst	0.114500	0.226700	0.382900	0.785500
concave points_worst	0.064930	0.099930	0.161400	0.291000
symmetry_worst	0.250400	0.282200	0.317900	0.419150
fractal_dimension_worst	0.071460	0.080040	0.092080	0.123010

```
[118]: df= df.drop(["Unnamed: 32","id"], axis = 1)
```

[120]: df

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	M	17.99	10.38	122.80	1001.0	
1	M	20.57	17.77	132.90	1326.0	
2	M	19.69	21.25	130.00	1203.0	
3	M	11.42	20.38	77.58	386.1	
4	M	20.29	14.34	135.10	1297.0	
..	...	...	...	...	...	
564	M	21.56	22.39	142.00	1479.0	
565	M	20.13	28.25	131.20	1261.0	
566	M	16.60	28.08	108.30	858.1	
567	M	20.60	29.33	140.10	1265.0	
568	B	7.76	24.54	47.92	181.0	
	smoothness_mean	compactness_mean	concavity_mean	concave_points_mean	points_mean	\
0	0.11840	0.27760	0.30010		0.14710	
1	0.08474	0.07864	0.08690		0.07017	
2	0.10960	0.15990	0.19740		0.12790	
3	0.14250	0.28390	0.24140		0.10520	
4	0.10030	0.13280	0.19800		0.10430	
..	...	...	...	...	...	
564	0.11100	0.11590	0.24390		0.13890	
565	0.09780	0.10340	0.14400		0.09791	
566	0.08455	0.10230	0.09251		0.05302	
567	0.11780	0.27700	0.35140		0.15200	
568	0.05263	0.04362	0.00000		0.00000	
	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	\
0	0.2419	...	25.380	17.33	184.60	
1	0.1812	...	24.990	23.41	158.80	
2	0.2069	...	23.570	25.53	152.50	
3	0.2597	...	14.910	26.50	98.87	
4	0.1809	...	22.540	16.67	152.20	
..	...	...	...	...	...	
564	0.1726	...	25.450	26.40	166.10	
565	0.1752	...	23.690	38.25	155.00	
566	0.1590	...	18.980	34.12	126.70	
567	0.2397	...	25.740	39.42	184.60	
568	0.1587	...	9.456	30.37	59.16	
	area_worst	smoothness_worst	compactness_worst	concavity_worst	\	
0	2019.0	0.16220	0.66560	0.7119		
1	1956.0	0.12380	0.18660	0.2416		
2	1709.0	0.14440	0.42450	0.4504		
3	567.7	0.20980	0.86630	0.6869		
4	1575.0	0.13740	0.20500	0.4000		

```

...
564      2027.0        0.14100        0.21130        0.4107
565      1731.0        0.11660        0.19220        0.3215
566      1124.0        0.11390        0.30940        0.3403
567      1821.0        0.16500        0.86810        0.9387
568      268.6         0.08996        0.06444        0.0000

    concave points_worst  symmetry_worst  fractal_dimension_worst
0            0.2654        0.4601        0.11890
1            0.1860        0.2750        0.08902
2            0.2430        0.3613        0.08758
3            0.2575        0.6638        0.17300
4            0.1625        0.2364        0.07678
...
564            ...          ...          ...
565            0.2216        0.2060        0.07115
566            0.1628        0.2572        0.06637
567            0.1418        0.2218        0.07820
568            0.2650        0.4087        0.12400
568            0.0000        0.2871        0.07039

```

[569 rows x 31 columns]

```
[122]: df["diagnosis"] = df["diagnosis"].astype("category")
```

```
[124]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   diagnosis        569 non-null    category
 1   radius_mean      569 non-null    float64
 2   texture_mean     569 non-null    float64
 3   perimeter_mean   569 non-null    float64
 4   area_mean        569 non-null    float64
 5   smoothness_mean  569 non-null    float64
 6   compactness_mean 569 non-null    float64
 7   concavity_mean   569 non-null    float64
 8   concave_points_mean 569 non-null    float64
 9   symmetry_mean    569 non-null    float64
 10  fractal_dimension_mean 569 non-null    float64
 11  radius_se        569 non-null    float64
 12  texture_se       569 non-null    float64
 13  perimeter_se     569 non-null    float64
 14  area_se          569 non-null    float64
 15  smoothness_se    569 non-null    float64
 16  compactness_se   569 non-null    float64

```

```

17 concavity_se           569 non-null    float64
18 concave_points_se      569 non-null    float64
19 symmetry_se            569 non-null    float64
20 fractal_dimension_se   569 non-null    float64
21 radius_worst           569 non-null    float64
22 texture_worst          569 non-null    float64
23 perimeter_worst        569 non-null    float64
24 area_worst              569 non-null    float64
25 smoothness_worst       569 non-null    float64
26 compactness_worst      569 non-null    float64
27 concavity_worst        569 non-null    float64
28 concave_points_worst   569 non-null    float64
29 symmetry_worst          569 non-null    float64
30 fractal_dimension_worst 569 non-null    float64
dtypes: category(1), float64(30)
memory usage: 134.2 KB

```

```
[126]: num_cols = df.select_dtypes(include=["float","int"]).columns

for col in num_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    upper_bound = Q3 + 1.5*IQR
    lower_bound = Q1 - 1.5*IQR
    df[col] = np.where(df[col] < lower_bound , lower_bound , df[col])
    df[col] = np.where(df[col] > upper_bound , upper_bound , df[col])
```

```
[128]: df["diagnosis"].value_counts()
```

```
[128]: diagnosis
B    357
M    212
Name: count, dtype: int64
```

```
[130]: y = df['diagnosis'].map({'B':0, 'M':1}) # benign = 0, malignant = 1
X = df.drop(columns=['diagnosis'])
```

```
[132]: # Identify numeric and categorical columns
num_cols = X.select_dtypes(include=['int64','float64']).columns.tolist()
cat_cols = X.select_dtypes(include=['object','category']).columns.tolist()
```

```
[134]: from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
```

```
[136]: num_transformer = Pipeline(steps=[('scaler', StandardScaler())])

cat_transformer = Pipeline(steps=[('OneHot', OneHotEncoder(handle_unknown = "ignore", sparse_output = False))])
```

```
preprocessor = ColumnTransformer(transformers=[('num', num_transformer, num_cols),
                                              ('cat', cat_transformer, cat_cols)],
                                              remainder='drop')
)
```

```
[138]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
print("Training set size:", X_train.shape)
print("Test set size:", X_test.shape)
```

```
Training set size: (455, 30)
Test set size: (114, 30)
```

```
[140]: models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(random_state=42),
    "XGBoost": XGBClassifier(eval_metric='logloss', random_state=42),
    "SVM": SVC(probability=True, random_state=42) # probability=True needed for ROC/AUC
}

for name, model in models.items():
    # Fit model
    model.fit(X_train, y_train)

    # Predictions
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:, 1] # needed for ROC/AUC

    print(f"\n== {name} ==")
    # Classification report
    print(classification_report(y_test, y_pred))

    # Confusion matrix
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(5, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f'{name} - Confusion Matrix')
```

```

plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# AUC-ROC
auc = roc_auc_score(y_test, y_prob)
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
plt.figure(figsize=(5,4))
plt.plot(fpr, tpr, label=f'AUC = {auc:.2f}')
plt.plot([0,1],[0,1], 'k--')
plt.title(f'{name} - ROC Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()

```

==== Logistic Regression ====

	precision	recall	f1-score	support
0	0.93	0.99	0.96	72
1	0.97	0.88	0.93	42
accuracy			0.95	114
macro avg	0.95	0.93	0.94	114
weighted avg	0.95	0.95	0.95	114

```

C:\Users\karan\anaconda3\Lib\site-
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

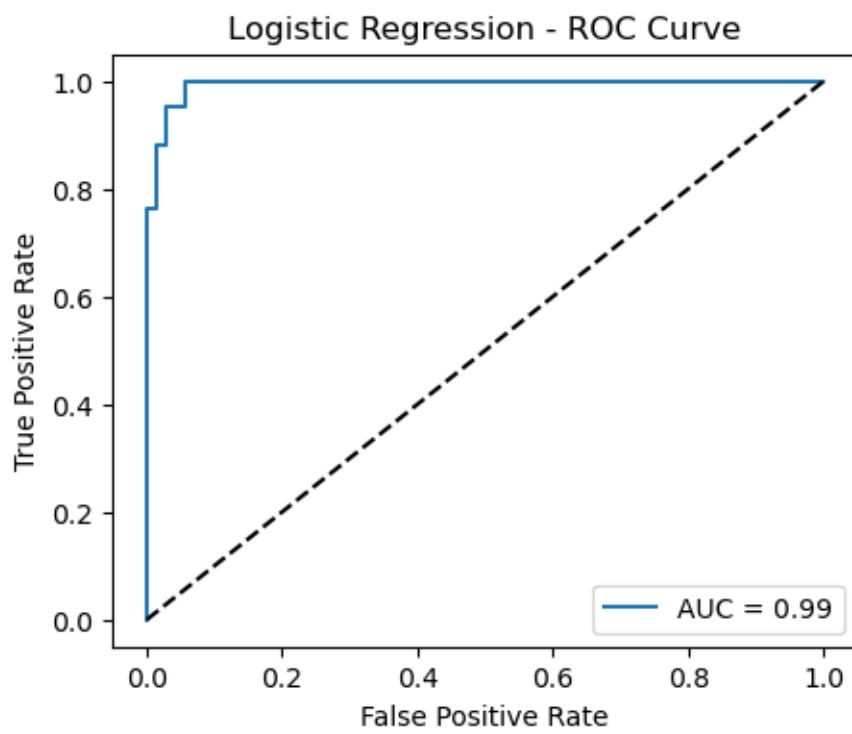
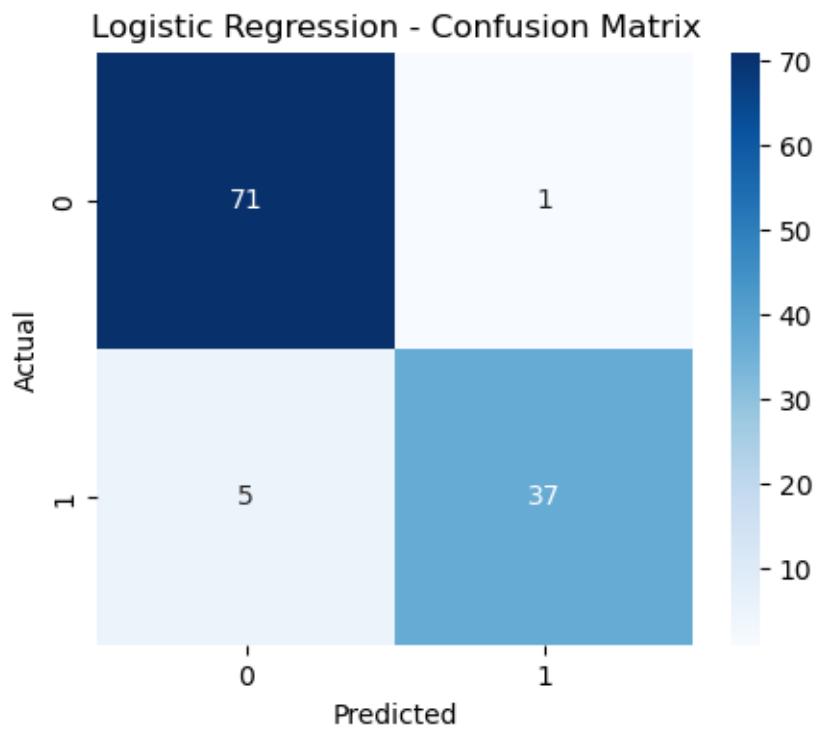
Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

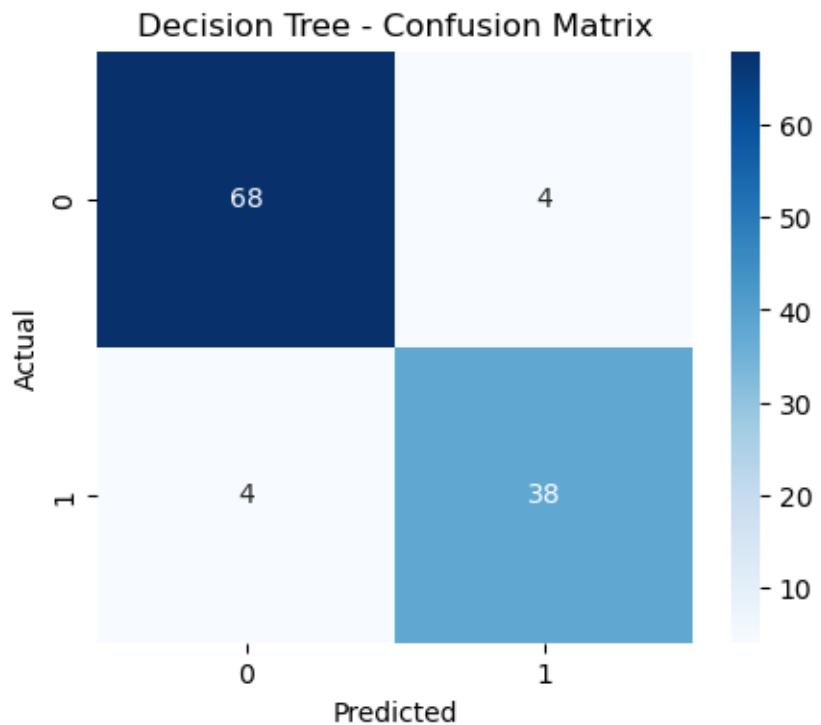
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

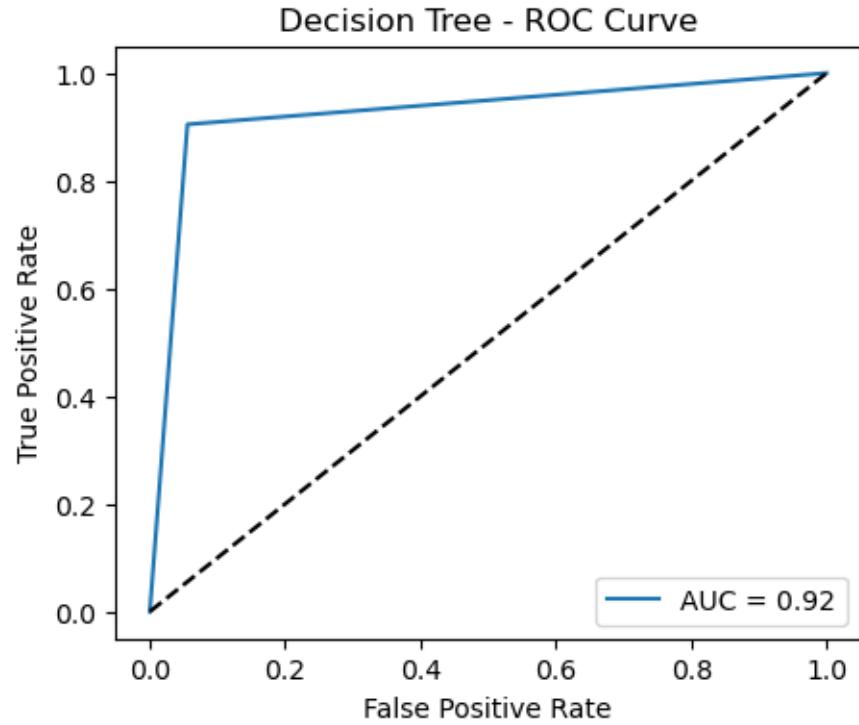
n\_iter\_i = \_check\_optimize\_result(



==== Decision Tree ===

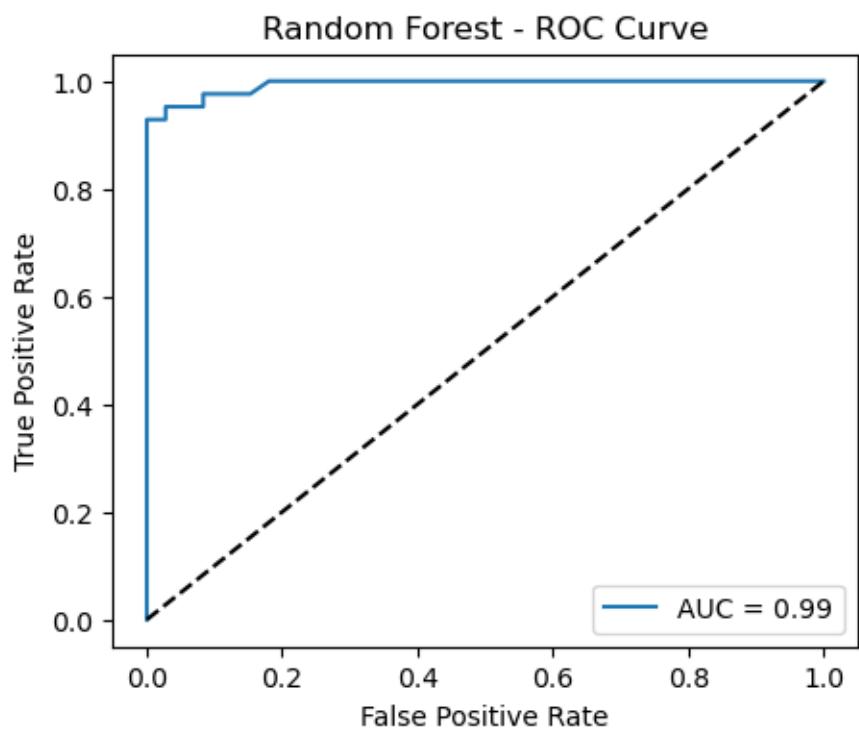
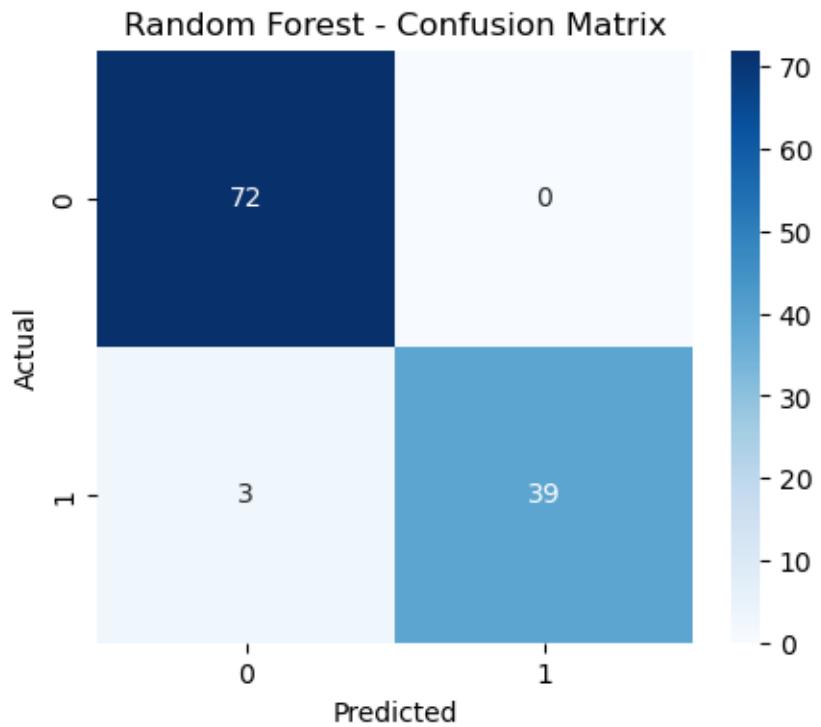
	precision	recall	f1-score	support
0	0.94	0.94	0.94	72
1	0.90	0.90	0.90	42
accuracy			0.93	114
macro avg	0.92	0.92	0.92	114
weighted avg	0.93	0.93	0.93	114





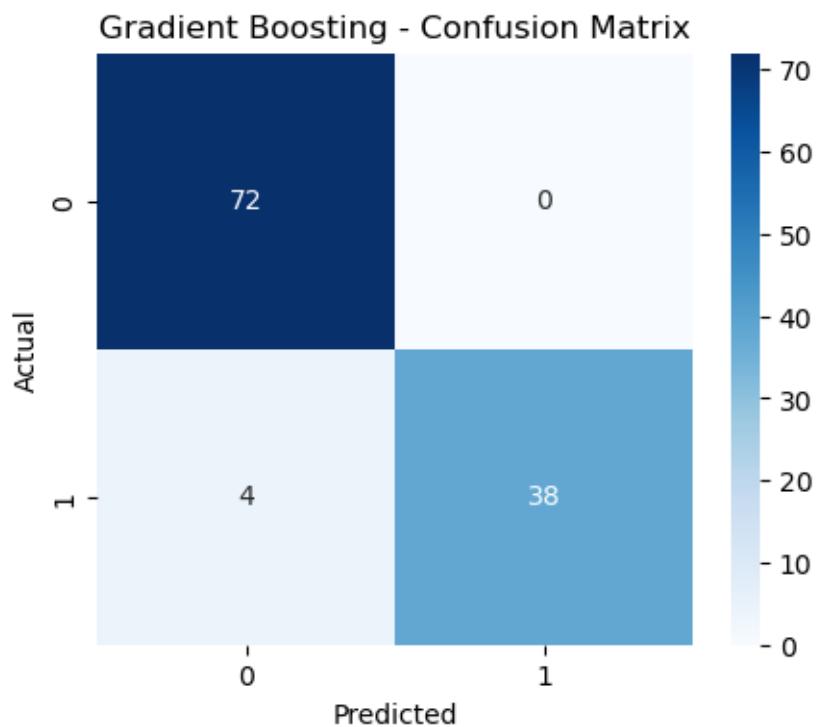
==== Random Forest ===

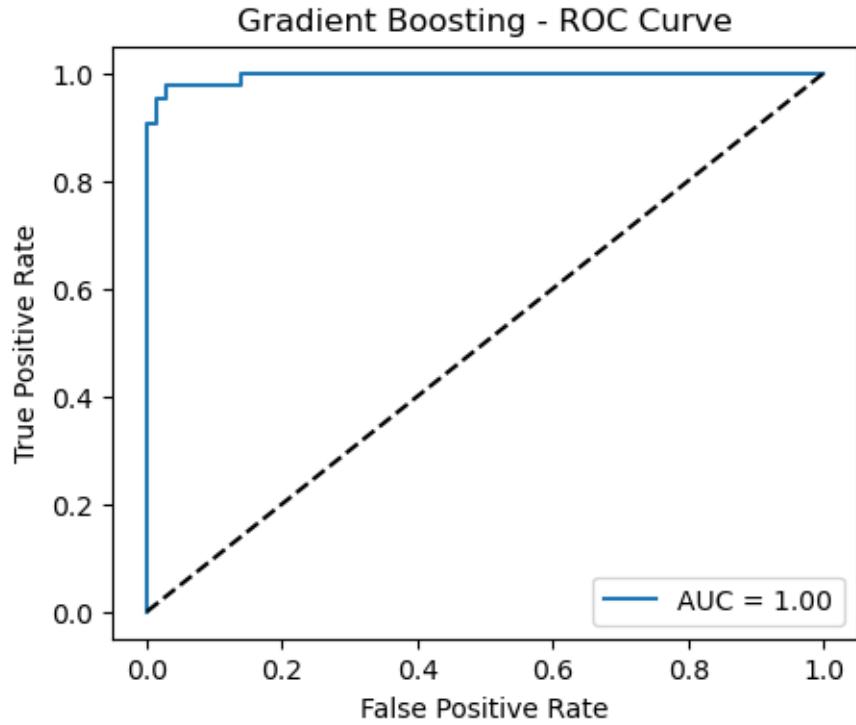
	precision	recall	f1-score	support
0	0.96	1.00	0.98	72
1	1.00	0.93	0.96	42
accuracy			0.97	114
macro avg	0.98	0.96	0.97	114
weighted avg	0.97	0.97	0.97	114



==== Gradient Boosting ===

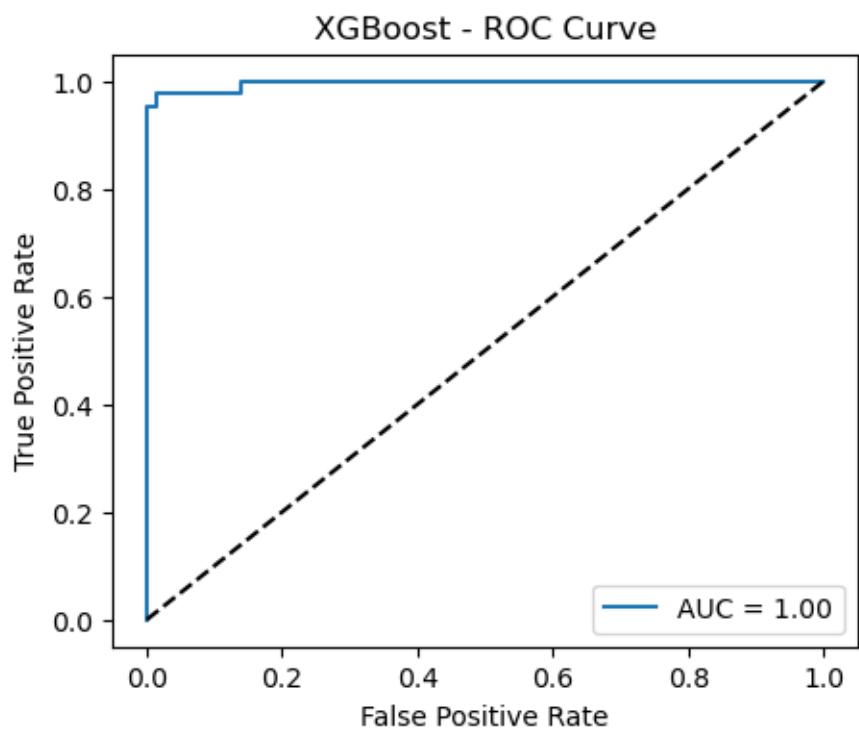
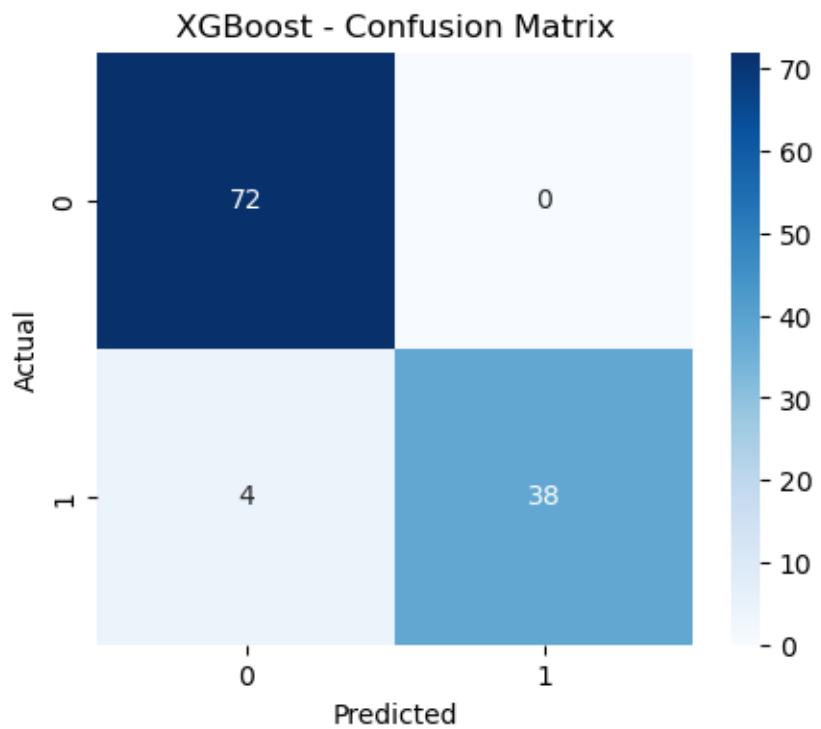
	precision	recall	f1-score	support
0	0.95	1.00	0.97	72
1	1.00	0.90	0.95	42
accuracy			0.96	114
macro avg	0.97	0.95	0.96	114
weighted avg	0.97	0.96	0.96	114





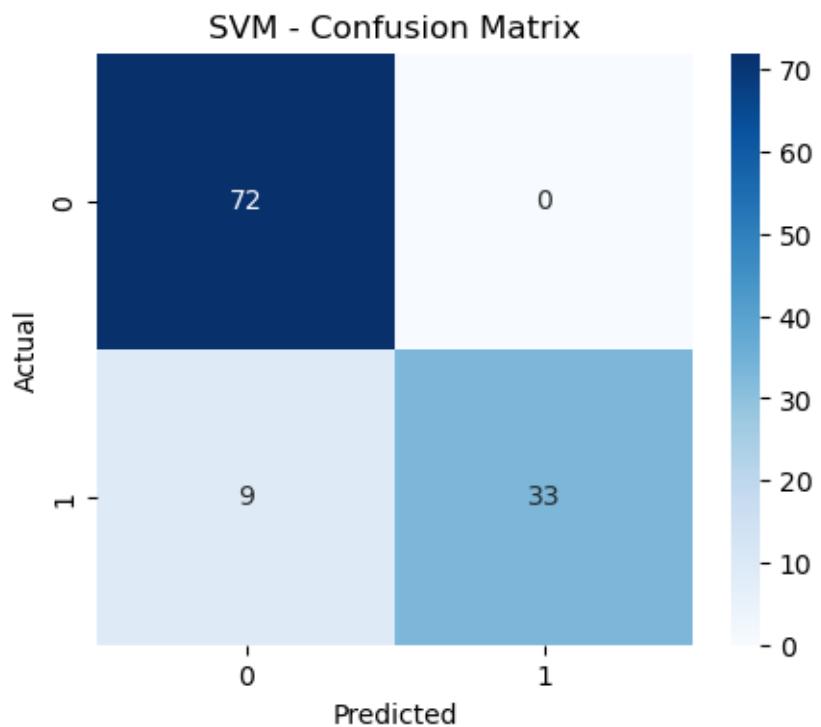
==== XGBoost ===

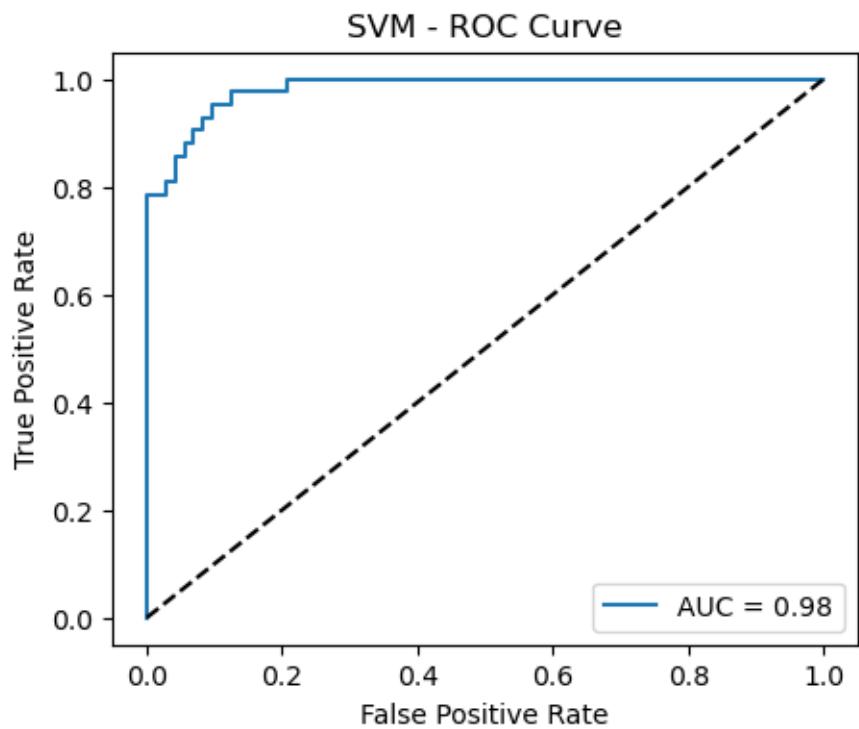
	precision	recall	f1-score	support
0	0.95	1.00	0.97	72
1	1.00	0.90	0.95	42
accuracy			0.96	114
macro avg	0.97	0.95	0.96	114
weighted avg	0.97	0.96	0.96	114



==== SVM ===

	precision	recall	f1-score	support
0	0.89	1.00	0.94	72
1	1.00	0.79	0.88	42
accuracy			0.92	114
macro avg	0.94	0.89	0.91	114
weighted avg	0.93	0.92	0.92	114





[102]:

[ ]: